

CSE 242 Assignment 4, Fall 2022

3 Questions, 100 pts, due: 23:59 pm, Nov 11th, 2022

Your name:Nistha Kumar

Student

ID:2005437

Instruction

- Submit your assignments onto **Gradescope** by the due date. Upload a **zip** file containing:

(1) The saved/latest **.ipynb** file, please **rename this file with your name included**.

(2) Also save your file into a pdf version, if error appears, save an html version instead (easy to grade for written questions).

For assignment related questions, please reach TA or grader through Slack/Piazza/Email.

- This is an **individual** assignment. All help from others (from the web, books other than text, or people other than the TA or instructor) must be clearly acknowledged.

Objective

- **Task 1:** EM algorithm (Mathematical Derivation)
- **Task 2:** K-Means implementation (Coding)
- **Task 3:** Kernel Methods with Noisy Setting (Coding)

Question 1. (EM algorithm, 20 pts)

Derive the E-step and M-step update equations of EM algorithm for estimating the Gaussian mixture model $p(X; \theta) = \sum_{k=1}^K \pi_k N(x; \mu_k, \sigma_k^2)$ where π_k is the mixture weight with $\pi_k \geq 0$ and $\sum_{k=1}^K \pi_k = 1$, and μ_k, σ_k^2 are the mean and variance of the gaussian distribution corresponding to cluster k.

For the E-step, first prove that $z_{ik} = P(z_i = k | X, \mu, \sigma, \pi) = \frac{\pi_k N(x_i; \mu_k, \sigma_k^2)}{\sum_{k=1}^K \pi_k N(x_i; \mu_k, \sigma_k^2)}$.

Then, for the M-step, show the derivation to compute the updates for (μ_k, π_k) . Note that, you don't need to show the derivation for σ_k . For each derivation step, mention the concept applied (e.g. just 2-3 keywords, e.g. formula for expectation, independence of datapoints, $(f+g)' = f' + g'$, etc ...).

Hint: For the M-step, you need to solve for

$$\mu_k^t = \underset{\mu_k}{\operatorname{argmax}} E_{p(Z|X, \mu^{(t-1)}, \sigma^{(t-1)}, \pi^{(t-1)})} [\log p(X, Z | \mu, \sigma, \pi)] \text{ (and similarly for } \pi_k \text{) by}$$

applying the first order conditions for function optimization (take derivative and set it to zero). Note that the term $p(Z|X, \mu^{(t-1)}, \sigma^{(t-1)}, \pi^{(t-1)})$ is the one computed in the E-step, and uses fixed values for μ, σ, π from the previous iteration (t-1).

```
In [1]: from IPython.display import Image
# Replace the figure name
Image(filename='1.1.jpg')
```

Out[1]:

1) Expectation Step:

The probability of assigning the i^{th} data point to the cluster k is given by:

$$z_{ik} = P(z_{i=k} | x = x_i, \mu_k, \sigma_k, \pi_k)$$

$$P(x|y) = \frac{P(y|x) P(x)}{P(y)} \quad \text{--- By Bayes rule}$$

$$\text{Then } z_{ik} = \frac{P(z_{i=k}) P(x_i, \mu_k, \sigma_k, \pi_k | z_{i=k})}{P(x_i, \mu_k, \sigma_k, \pi_k)}$$

$$\text{Since } P(z_{i=k}) = \pi_k \quad \text{--- (1)}$$

N -data points

$$\text{and } P(x_i, \mu_k, \sigma_k, \pi_k | z_{i=k}) = N(x_i | \mu_k, \sigma_k^2) \quad \text{--- (2)}$$

- likelihood of
datapoint given the
cluster k

$$P(x) = \sum_y P(x, y) = \sum_y P(x|y) P(y)$$

$$P(x_i, \mu_k, \sigma_k, \pi_k) = \sum_{k=1}^K P(x_i, \mu_k, \sigma_k, \pi_k | z_{i=k}) * P(z_{i=k})$$

$$\Rightarrow P(x_i, \mu_k, \sigma_k, \pi_k) = \sum_{k=1}^K N(x_i | \mu_k, \sigma_k^2) * \pi_k \quad \text{--- (3)}$$

$$\text{Hence : } z_{ik} = \frac{\pi_k * N(x_i | \mu_k, \sigma_k^2)}{\sum_{k=1}^K N(x_i | \mu_k, \sigma_k^2)} \quad \text{from (1) (2) \& (3)}$$

In [2]: `Image(filename='1.2.jpg')`

Out [2]:

Maximization Step

$$\pi^{\text{new}}, \mu^{\text{new}}, \sigma^{\text{new}} = \underset{\pi, \mu, \sigma}{\operatorname{argmax}} \mathbb{E}_{P(z|x, \pi^{t-1}, \mu^{t-1}, \sigma^{t-1})} \log(P(x, z | \pi, \mu, \sigma)) \quad \text{--- (1)}$$

$$\text{Assuming } \Omega = \mathbb{E}_{P(z|x, \pi^{t-1}, \mu^{t-1}, \sigma^{t-1})} \log(P(x, z | \pi, \mu, \sigma))$$

$$\text{Since } E[f(x)] = \sum_x f(x) P(x)$$

$$\text{Hence } \Omega = \sum_z \log(P(x, z | \pi, \mu, \sigma)) P(z | x, \mu^{t-1}, \sigma^{t-1}, \pi^{t-1})$$

$$\Omega = \sum_{z_1 \text{ to } z_N} P(z_1, z_2, \dots, z_N | x_1, x_2, \dots, x_N, \mu^{t-1}, \sigma^{t-1}, \pi^{t-1}) \log(P(z, x | \pi, \mu, \sigma))$$

$$= \sum_{z_1 \text{ to } N} P(z_1, z_2, \dots, z_N | x_1, x_2, \dots, x_N, \mu^{t-1}, \sigma^{t-1}, \pi^{t-1}) \times \log(P(z | \pi) \times P(x | \sigma, \mu, z))$$

If we consider data points to be independent

$$\Omega = \sum_{z_1 \text{ to } N} P(z_1, \dots, z_N | x_1, \dots, x_N, \mu^{t-1}, \sigma^{t-1}, \pi^{t-1}) \log \left[\prod_{i=1}^N P(z_i | \pi) \times P(x_i | z_i, \mu, \sigma) \right]$$

$$= \sum_{z_1 \text{ to } N} P(z_1, \dots, z_N | x_1, \dots, x_N, \mu^{t-1}, \sigma^{t-1}, \pi^{t-1}) \left[\sum_{i=1}^N \log \left[P(z_i | \pi) \times P(x_i | z_i, \mu, \sigma) \right] \right]$$

In [3]: `Image(filename='1.3.jpg')`

Out [3]:

$$\Omega = \sum_{i=1}^N \sum_{z_i \in \mathcal{N}} P(z_1, \dots, z_N | x_1, \dots, x_N, \mu^{t-1}, \sigma^{t-1}, \pi^{t-1}) \times \log \left[\frac{P(z_i | \pi) \times P(x_i | z_i, \mu, \sigma)}{P(x_i | z_i, \mu, \sigma)} \right]$$

z_1, z_2, \dots, z_N are all independent

Plus: $\Omega =$

$$\sum_{i=1}^N \sum_{z_i \in \mathcal{N}} P(z_1 | x_1, \mu^{t-1}, \sigma^{t-1}, \pi^{t-1}) \times P(z_2 | x_2, \mu^{t-1}, \sigma^{t-1}, \pi^{t-1}) \times \dots \times P(z_N | x_N, \mu^{t-1}, \sigma^{t-1}, \pi^{t-1}) \times \log [P(z_i | \pi) \times P(x_i | z_i, \mu, \sigma)]$$

$$\Omega = \sum_{i=1}^N \sum_{z_i} P(z_i | x_i, \mu, \sigma, \pi) \times \log [P(z_i | \pi) \cdot P(x_i | z_i, \mu, \sigma)] \times \underbrace{\sum_{z_1 \in \mathcal{N}} P(z_1 | x_1, \mu, \sigma, \pi) \cdot P(z_2 | x_2, \mu, \sigma, \pi) \cdot \dots \cdot P(z_N | x_N, \mu, \sigma, \pi)}_{=1}$$

$$\Omega = \sum_{i=1}^N \sum_k P(z_i | x_i, \mu, \sigma, \pi^{t-1}) \times \log [P(z_i | \pi) \cdot P(x_i | z_i, \mu, \sigma)] \quad (2)$$

$$\mu^{\text{new}} = \operatorname{argmax}_{\mu} \Omega(\mu)$$

$$= \operatorname{argmax}_{\mu} \sum_{i=1}^N \sum_{k'} P(z_i | x_i, \mu^{t-1}) \times \log [P(z_i = k' | \pi) \cdot P(x_i | z_i = k', \mu, \sigma, z_i)]$$

$$\frac{\partial \Omega(\mu)}{\partial \mu_k} = 0$$

In [4]: `Image(filename='1.4.jpg')`

Out [4]:

$$\text{Thus } \frac{\partial Q(\mu)}{\partial \mu_k} = \frac{\partial}{\partial \mu_k} \left[\sum_{i=1}^N \sum_{k'} P(z_i = k' | x_i, \mu^{t-1}) \times \log(P(z_i = k' | \pi) \cdot P(x_i | \mu, \sigma, z_i)) \right] = 0$$

Derivative of sum is the sum of derivatives

$$\sum_{i=1}^N \sum_{k'} P(z_i = k' | x_i, \mu^{t-1}) \frac{\partial}{\partial \mu_k} \underbrace{\log(P(z_i = k' | \pi) \cdot P(x_i | \mu, \sigma, z_i))}_{N(x_i | \mu_{k'}, \sigma_{k'}, \pi_{k'})}$$

$$\Rightarrow \sum_{i=1}^N P(z_i = k | x_i, \mu^{t-1}) \frac{\partial}{\partial \mu_k} [\log(N(x_i | \mu_{k'}, \sigma_{k'}))] = 0$$

$$\Rightarrow \sum_{i=1}^N P(z_i = k | x_i, \mu^{t-1}) \frac{\partial}{\partial \mu_k} \left[\log\left(\frac{1}{\sqrt{2\pi}\sigma_k} \times e^{-\frac{(x_i - \mu_k)^2}{2\sigma^2}}\right) \right] = 0$$

$$\sum_{i=1}^N P(z_i = k | x_i, \mu^{t-1}) \frac{\partial}{\partial \mu_k} \left[\log\left(\frac{1}{\sqrt{2\pi}\sigma_k}\right) - \frac{(x_i - \mu_k)^2}{2\sigma^2} \times \log e \right] = 0$$

$$\sum_{i=1}^N P(z_i = k | x_i, \mu^{t-1}) \frac{\partial}{\partial \mu_k} \left[-\frac{(x_i - \mu_k)^2}{2\sigma^2} \right] = 0$$

$$\sum_{i=1}^N z_{ik} \left(\frac{-2 \times (x_i - \mu_k)}{2\sigma^2} \times -1 \right) = 0$$

$$\sum_{i=1}^N z_{ik} (x_i - \mu_k) = 0$$

In [5]: `Image(filename='1.5.jpg')`

Out [5]:

$$\Rightarrow \sum_{i=1}^N z_{ik} x_i - \sum_{i=1}^N z_{ik} \mu_k = 0$$

$$\text{or } \mu_k = \frac{\sum_{i=1}^N z_{ik} x_i}{\sum_{i=1}^N z_{ik}}$$

$$\pi_k^{\text{new}} = \text{argmax} \mathbb{E}_{P(z|x, \pi^{t-1}, \mu^{t-1}, \sigma^{t-1})} \ln(P(x, z | \pi, \mu, \sigma)) - (3)$$

From Lagrange's condition

$$L(\pi, \lambda) = \Omega(\pi) + \lambda \left(1 - \sum_{k'=1}^K \pi_{k'} \right) - (4)$$

$$\frac{\partial}{\partial \pi_k} L(\pi, \lambda) = 0$$

$$\frac{\partial}{\partial \pi_k} \Omega(\pi) + \frac{\partial}{\partial \pi_k} \left(\lambda \left(1 - \sum_{k'=1}^K \pi_{k'} \right) \right) = 0$$

$$\text{From (3)} \quad \frac{\partial}{\partial \pi_k} \left(\mathbb{E}_{P(z|x, \pi, \mu, \sigma)} \ln(P(z, x | \pi, \mu, \sigma)) \right) + \frac{\partial}{\partial \pi_k} \left(\lambda \left(1 - \sum_{k'=1}^K \pi_{k'} \right) \right) = 0$$

In [6]: `Image(filename='1.6.jpg')`

Out [6]:

$$\frac{\partial}{\partial \pi_k} \left[\sum_{i=1}^N P(z_i=k | x_i, \pi^{t-1}) \log(P(z_i | \pi) P(x_i | \mu, \sigma, z_i)) \right] + \lambda \frac{\partial}{\partial \pi_k} \left[1 - \sum_{k'=1}^K \pi_{k'} \right] = 0$$

$$\text{or } \frac{\partial}{\partial \pi_k} \left[\sum_{i=1}^N z_{ik} (\log[\pi_k \times N(x_i | \mu, \sigma, z_i)]) \right] - \lambda \frac{\partial}{\partial \pi_k} [1 + \dots + k] = 0$$

$$\frac{\partial}{\partial \pi_k} \left[\sum_{i=1}^N z_{ik} (\log \pi_k + \log(N(x_i | \mu, \sigma, z_i))) \right] - \lambda = 0$$

$$\sum_{i=1}^N z_{ik} \frac{\partial}{\partial \pi_k} (\log \pi_k + \log(N(x_i | \mu, \sigma, z_i))) = \lambda$$

$$\Rightarrow \sum_{i=1}^N z_{ik} \left(\frac{1}{\pi_k} + 0 \right) = \lambda$$

$$\text{or } \boxed{\lambda = \frac{\sum_{i=1}^N z_{ik}}{\pi_k}} \quad \text{--- (5)}$$

$$\text{We know that: } \sum_{k'=1}^K P(z_i=k') = \sum_{k'=1}^K \pi_{k'} = 1$$

$$\text{Thus } \lambda = \lambda \times 1 \quad \text{or} \quad \lambda = \sum_{i=1}^K \lambda \times \pi_{k'}$$

In [7]: `Image(filename='1.7.jpeg')`

Out[7]:

Substituting the value of λ from (5)

$$\lambda = \sum_{k'=1}^K \left(\frac{\sum_{i=1}^N z_{ik'}}{\pi_{k'}} \right) \pi_{k'}$$

$$\lambda = \sum_{i=1}^N \sum_{k'=1}^K z_{ik'}$$

$\therefore \sum_{k'=1}^K z_{ik'} = 1$ because we are

$$\therefore \lambda = \sum_{i=1}^N 1 = N \quad \text{--- (6)}$$

adding all the probabilities of the datapoint belonging to all the clusters.

from (5) & (6) :

$$\pi_k = \frac{\sum_{i=1}^N z_{ik}}{N}$$

References :

<https://www.youtube.com/watch?v=XGRXH9Fcc-w>

Question 2. (K-Means implementation, 20 pts)

Question 2.1. Implement K-means in Python from scratch. Complete following sub-functions `update_centroids` and `update_assignments`.

```

In [8]: import numpy as np
from scipy.spatial.distance import cdist
def update_assignments(data, centroids):

    #####
    #### YOUR CODE HERE ####
    row, col = data.shape
    assignments = np.empty([row])
    distances = cdist(data, centroids)
    assignments=np.argmin(distances,axis=1)

    ## you will get cluster#
    ##assignments here #####
    #####

    return assignments

def update_centroids(data,centroids,assignments):

    #####
    #### YOUR CODE HERE ####
    #####
    K = centroids.shape[0]
    centroids=np.empty(centroids.shape)
    for i in range(K):
        centroids[i]=np.mean(data[assignments ==i], axis=0)

    return centroids

def kmeans(data, centroids, max_iterations):

    for j in range(max_iterations):
        # update cluter assignments
        assignments = update_assignments(data,centroids)    # WRITE CODE FOR

        # update centroid locations
        centroids = update_centroids(data,centroids,assignments)    # WRITE CC

    # final assignment update
    assignments = update_assignments(data,centroids)
    return centroids, assignments

```

Question 2.2. Run your code on following toy dataset for different k-values, where $k = \{2, 3, 4, 6, 10\}$ and plot the cluster assignments for different k's as shown in following diagram.


```
In [9]: from sklearn.datasets import make_blobs
import matplotlib.pyplot as plt

# Generate sample data
n_samples = 4000
n_components = 4

X, y_true = make_blobs(
    n_samples=n_samples, centers=n_components, cluster_std=0.60, random_stat
)

colors = ["#4EACCC", "#FF9C34", "#4E9A06", "m"]

for k, col in enumerate(colors):
    cluster_data = y_true == k
    plt.scatter(X[cluster_data, 0], X[cluster_data, 1], c=col, marker=".", s
```



```



In [10]: import numpy as np




# function to get initial cluster centroids; we randomly choose k points from
def get_initial_clusters(k, X):
    random_indices = np.random.randint(0, X.shape[0], k)
    initial_centroids = X[random_indices]


    return initial_centroids

# your code here.
clusters=[2,3,4,6,10]
colors = ["#4EACC5", "#FF9C34", "#4E9A06", "m",
          "#fbc068", "#fd97e2", "#b8c2aa", "#7c5858",
          "#3b6f05", "#001b58"]
fig, axs = plt.subplots(len(clusters), figsize=[8,30])
for idx, i in enumerate(clusters):
    centroids=get_initial_clusters(i, X)
    centroids,assignments=kmeans(X,centroids,100)
    #print(assignments.shape)
    for k in range(i):
        cluster_data = assignments == k
        axs[idx].scatter(X[cluster_data, 0], X[cluster_data, 1], color=colors

```







Question 3. (Kernel Methods with Noisy Setting, 60 pts)

SVM on synthetic dataset generated as follows:

- Draw 1000 (x_0, x_1) feature vectors from the 2-D Gaussian distribution with mean $\mu_+ = (1, 1)$ and $\Sigma_+ = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}$ and label them as $+1$.
- Draw 1000 (x_0, x_1) feature vectors from the 2-D Gaussian distribution with mean $\mu_- = (-1, -1)$ and $\Sigma_- = \begin{bmatrix} 3 & 0 \\ 0 & 3 \end{bmatrix}$ and label them as -1 .
- This gives you a 2000 example training set. Repeat the above to draw a test set the same way.

Use a SVM package (`scikit-learn svm.SVC` class) to learn SVMs with a variety of parameter settings.

```

In [11]: # Your code here
import numpy as np
import pandas as pd
from sklearn.utils import shuffle
np.random.seed(220)
mean1 = (1,1)
cov1 = [[1,0],[0,1]]
mean2 = (-1,-1)
cov2 = [[3,0],[0,3]]
df1 = pd.DataFrame(np.random.multivariate_normal(mean1,cov1,1000),columns=['x','y'])
df1['label'] = +1
df2 = pd.DataFrame(np.random.multivariate_normal(mean2,cov2,1000),columns=['x','y'])
df2['label'] = -1
df_train = pd.concat([df1,df2], ignore_index=True)
#Shuffling to avoid consistency
# df_train = shuffle(df_train)
df_train = df_train.reset_index(drop=True)
df_train

```

```

Out[11]:

```

	x	y	label
0	-0.324352	-0.465870	1
1	-0.249991	1.654616	1
2	2.235651	-0.122390	1
3	1.414549	0.806798	1
4	-0.747681	2.687990	1
...
1995	-2.454008	-3.098631	-1
1996	-5.584851	-0.248678	-1
1997	2.095982	0.153912	-1
1998	1.639896	-3.109557	-1
1999	0.186215	-2.850072	-1

2000 rows × 3 columns

```

In [12]: df1 = pd.DataFrame(np.random.multivariate_normal(mean1,cov1,1000),columns=['x','y'])
df1['label'] = +1
df2 = pd.DataFrame(np.random.multivariate_normal(mean2,cov2,1000),columns=['x','y'])
df2['label'] = -1
df_test = pd.concat([df1,df2], ignore_index=True)
df_test

```


Out[12]:

	x	y	label
0	-1.898160	1.556240	1
1	1.784113	-0.436965	1
2	1.986831	2.236635	1
3	-0.213913	1.240479	1
4	2.431783	0.760874	1
...
1995	-1.012130	0.737504	-1
1996	0.199889	1.865863	-1
1997	-2.578212	1.371016	-1
1998	-0.800268	-1.951012	-1
1999	-2.646145	-1.403640	-1

2000 rows × 3 columns

(a -- 20 pts)

- Use an RBF kernel with parameters $C = 1$, $\gamma = 0.01$.
- For each training data with +1 label, randomly flip their label to -1 with probability **0.35**.
- For each training data with -1 label, randomly flip their label to +1 with probability **0.20**.
- Train with the above noisy training examples.
- Random flipping introduces the randomness. You can repeat multiple times (e.g. 20) and then report the average accuracy on the testing dataset (clean) in the noise parameter setting.

```
In [13]: from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.svm import SVC
import random
X_train = df_train[['x', 'y']]
y_train = df_train['label'].values
X_test = df_test[['x', 'y']]
y_test = df_test['label'].values

model = SVC(kernel='rbf', C=1, gamma=0.01, random_state = 21)
model = model.fit(X_train, y_train)

y_pred = model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy score before flip: ", accuracy)

Accuracy score before flip: 0.8585
```

```
In [14]: def flip_labels(pos, neg, prob_pos=0.35, prob_neg=0.20):
    pos_prob=np.random.random(size=pos.shape[0])
    flip_pos=np.where(pos_prob<prob_pos,-1,pos)

    neg_prob=np.random.random(size=neg.shape[0])
    flip_neg=np.where(neg_prob<prob_neg,1,neg)

    return np.hstack([flip_pos, flip_neg])

y_train_flipped= flip_labels(y_train[:1000], y_train[1000:])
```

```
In [15]: model1 = SVC(kernel='rbf', probability=True, C=1, gamma=0.01)
model1 = model1.fit(X_train, y_train_flipped)

y_pred = model1.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy score after flip: ", accuracy)

Accuracy score after flip: 0.7895
```

(b -- 20 pts) Open question

- Try using **K-Nearest Neighbors** to correct wrong labels before training.
- Then train the model with the newly processed training dataset.
- Report the accuracy on the testing dataset in the noise parameter setting. Do you observe performance improvement?

```

In [16]: # Your code here
import matplotlib.pyplot as plt
%matplotlib inline
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score
best_knn_svm=0
best_k_svm=0

best_knn=0
best_k=0
best_knn_model=None

for k in range(1,100):
    knn_neigh=KNeighborsClassifier(n_neighbors=k)
    knn_neigh.fit(X_train,y_train_flipped)

    y_neigh_predict_test=knn_neigh.predict(X_train)
    if(accuracy_score(y_test,y_neigh_predict_test)*100>best_knn):
        best_knn=accuracy_score(y_test,y_neigh_predict_test)*100
        best_k=k
        best_knn_model=knn_neigh

#print(best_k)
svm_knn=SVC(kernel='rbf',random_state=0,gamma=.01,C=1)

y_train_knn=best_knn_model.predict(X_train)
svm_knn.fit(X_train,y_train_knn)

y_svm_knn_predict_test=svm_knn.predict(X_test)

best_knn_svm=accuracy_score(y_test,y_svm_knn_predict_test)*100

    #if(accuracy_score(y_test,y_svm_knn_predict_test)*100>best_knn_svm):
    #    best_knn_svm=accuracy_score(y_test,y_svm_knn_predict_test)*100
    #    best_k_svm=k

print(f"Accuracy with knn+svm: {best_knn_svm}")
#print(f"Accuracy with just knn: {best_knn},k neighbors:{best_k}")

```

Accuracy with knn+svm: 83.89999999999999

After correcting the noisy labels using knn, and then running svm on that- the accuracy has improved

(c -- 20 pts) Open question

- Try using **clustering (i.e., K-means, EM-clustering)** to correct wrong labels before training.
- Then train the model with the newly processed training dataset.
- Report the accuracy on the testing dataset in the noise parameter setting. Do you observe performance improvement?

```
In [17]: # Your code here
from sklearn.cluster import KMeans
kmeans_cluster=KMeans(n_clusters=2,random_state=0).fit(X_train)
Y_kmeans=kmeans_cluster.labels_
#print(Y_kmeans)
for i,ele in enumerate(Y_kmeans):
    if Y_kmeans[i] ==1:
        Y_kmeans[i]=-1
    elif Y_kmeans[i]==0:
        Y_kmeans[i]=1
#print(Y_kmeans)
svm_kmeans=SVC(kernel='rbf', random_state=0,gamma=.01,C=1)
svm_kmeans.fit(X_train,Y_kmeans)

y_svm_kmeans_predict_test=svm_kmeans.predict(X_train)

print(f"Accuracy with kmeans+svm: {accuracy_score(y_test,y_svm_kmeans_predic

Accuracy with kmeans+svm: 82.25
```

The accuracy after k-means is comparatively less as it is unsupervised and the labels are labelled based on the cluster. But the accuracy has still improved a little.

```
In [18]: from sklearn.mixture import GaussianMixture
EM_cluster= GaussianMixture(n_components=2,random_state=0).fit(X_train)
Y_EM=EM_cluster.predict(X_train)
#print(Y_EM)

for i,ele in enumerate(Y_EM):
    if Y_EM[i] ==1:
        Y_EM[i]==-1
    elif Y_EM[i]==0:
        Y_EM[i]=1
#print(Y_EM)
svm_EM=SVC(kernel='rbf', random_state=0,gamma=.01,C=1)
svm_EM.fit(X_train,Y_EM)

y_svm_EM_predict_test=svm_EM.predict(X_train)

print(f"Accuracy with EM+svm: {accuracy_score(y_test,y_svm_EM_predict_test)}")
Accuracy with EM+svm: 83.35000000000001
```

The accuracy has improved compared to the one with noisy labels. Still its less than knn because EM is again unsupervised.